

Diabetes Predictions By using Machine Learning

IMPORT LIBRARIES AS WELL AS DATASET

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
import os
%matplotlib inline
warnings.filterwarnings('ignore')
```


```
In [2]: df = pd.read_csv("Downloads/diabetes_prediction_.csv")
```

```
In [3]: df
```

Out[3]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood
0	Female	80.0	0	1	never	25.19	6.6	
1	Female	54.0	0	0	No Info	27.32	6.6	
2	Male	28.0	0	0	never	27.32	5.7	
3	Female	36.0	0	0	current	23.45	5.0	
4	Male	76.0	1	1	current	20.14	4.8	
...
99995	Female	80.0	0	0	No Info	27.32	6.2	
99996	Female	2.0	0	0	No Info	17.37	6.5	
99997	Male	66.0	0	0	former	27.83	5.7	
99998	Female	24.0	0	0	never	35.42	4.0	
99999	Female	57.0	0	0	current	22.43	6.6	

100000 rows × 9 columns



We make two copy of the original dataset

```
In [4]: df1 = df.copy(deep=True)
df2 = df.copy(deep=True)
```

DATA PREPROCESSING

In [5]: `df.isnull().sum()`

```
Out[5]: gender          0
age                  0
hypertension        0
heart_disease       0
smoking_history     0
bmi                 0
HbA1c_level         0
blood_glucose_level 0
diabetes            0
dtype: int64
```

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                100000 non-null  object
1   age                   100000 non-null  float64
2   hypertension          100000 non-null  int64
3   heart_disease         100000 non-null  int64
4   smoking_history       100000 non-null  object
5   bmi                   100000 non-null  float64
6   HbA1c_level          100000 non-null  float64
7   blood_glucose_level   100000 non-null  int64
8   diabetes              100000 non-null  int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

In [7]: `df.describe()`

```
Out[7]:
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_gluc
count	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000	1000
mean	41.885856	0.07485	0.039420	27.320767	5.527507	1
std	22.516840	0.26315	0.194593	6.636783	1.070672	
min	0.080000	0.00000	0.000000	10.010000	3.500000	
25%	24.000000	0.00000	0.000000	23.630000	4.800000	1
50%	43.000000	0.00000	0.000000	27.320000	5.800000	1
75%	60.000000	0.00000	0.000000	29.580000	6.200000	1
max	80.000000	1.00000	1.000000	95.690000	9.000000	3

In [8]: `df.columns`



```
Out[8]: Index(['gender', 'age', 'hypertension', 'heart_disease', 'smoking_histor
y',
              'bmi', 'HbA1c_level', 'blood_glucose_level', 'diabetes'],
              dtype='object')
```


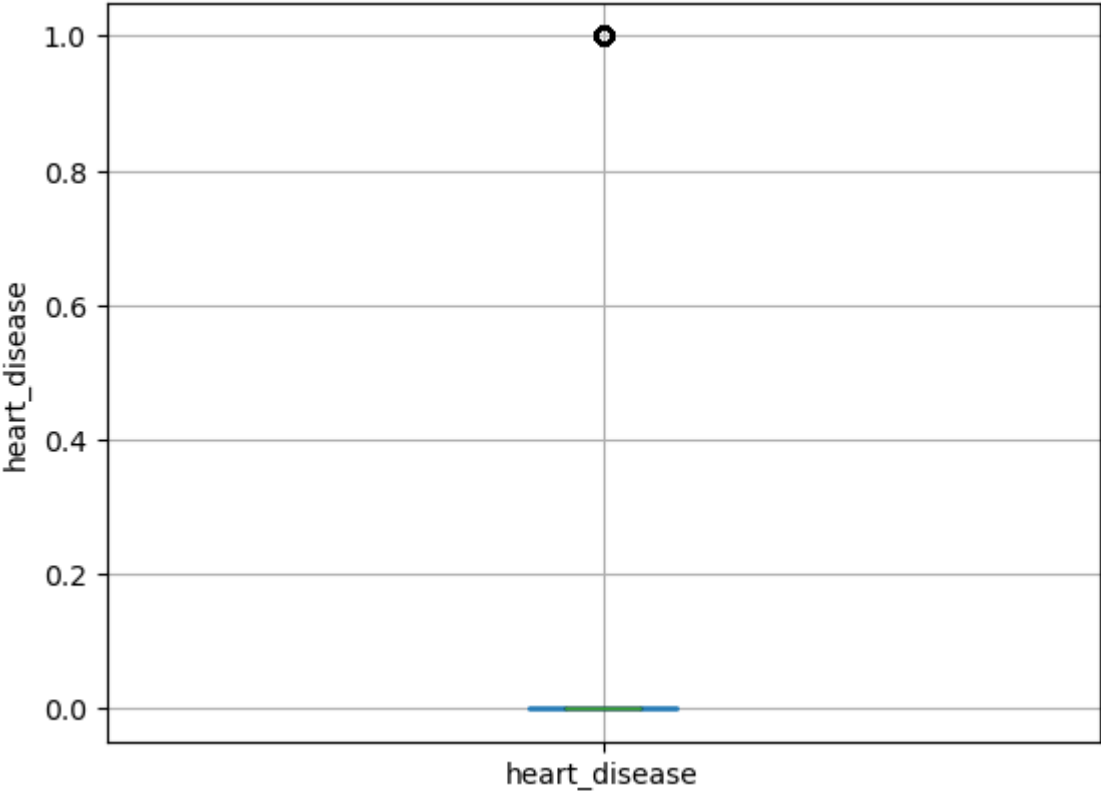
EXPLORATORY DATA ANALYSIS

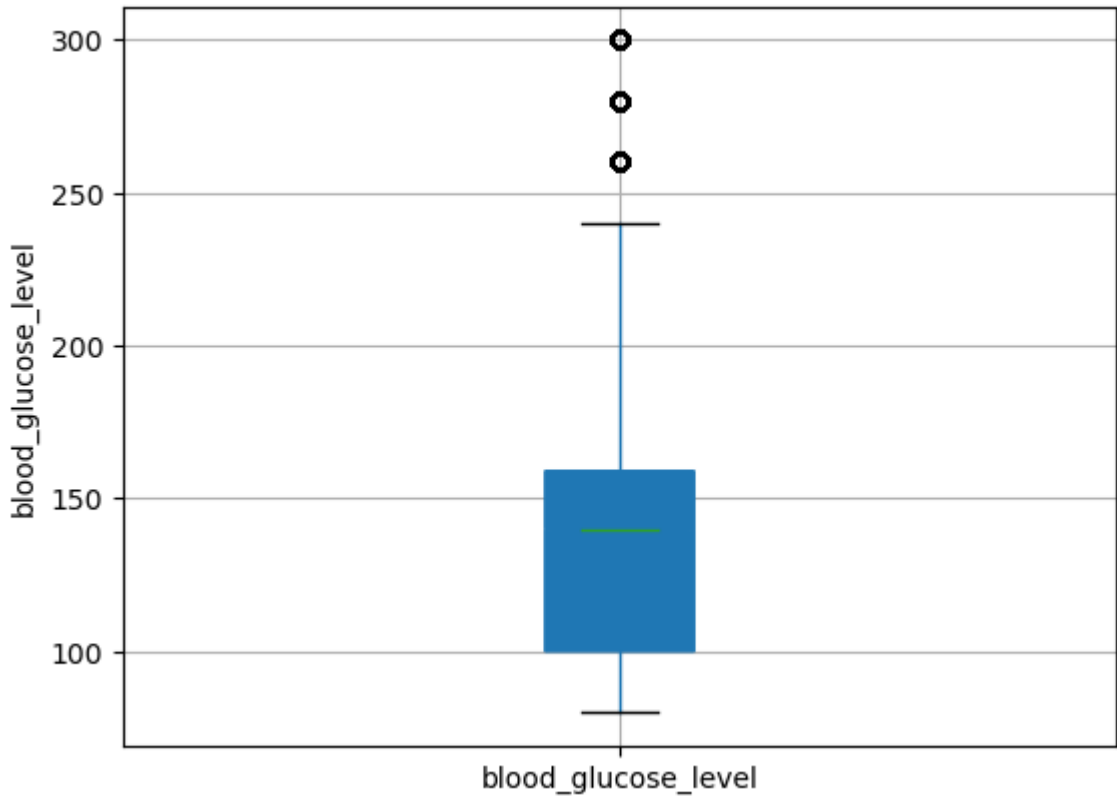

Here, in this dataset have Categorical features for that we separate categorical and numerical features. So, now create EDA for numerical data. First, check the Outliers


```
In [9]: cat = df.select_dtypes(include="object")  
        num = df.select_dtypes(include="number")
```

```
In [10]: for i in num:
          num.boxplot(column=i, patch_artist = True, notch = 'True')
          plt.ylabel(i)
          plt.show()
```

















```
In [11]: for i in num:
          sns.distplot(df[i], kde = True, color = 'green')
          plt.show()
```













Now, we Check the categorical distrubition.



```
In [12]: for i in cat:  
         sns.countplot(df, x = df[i])  
         plt.show()
```


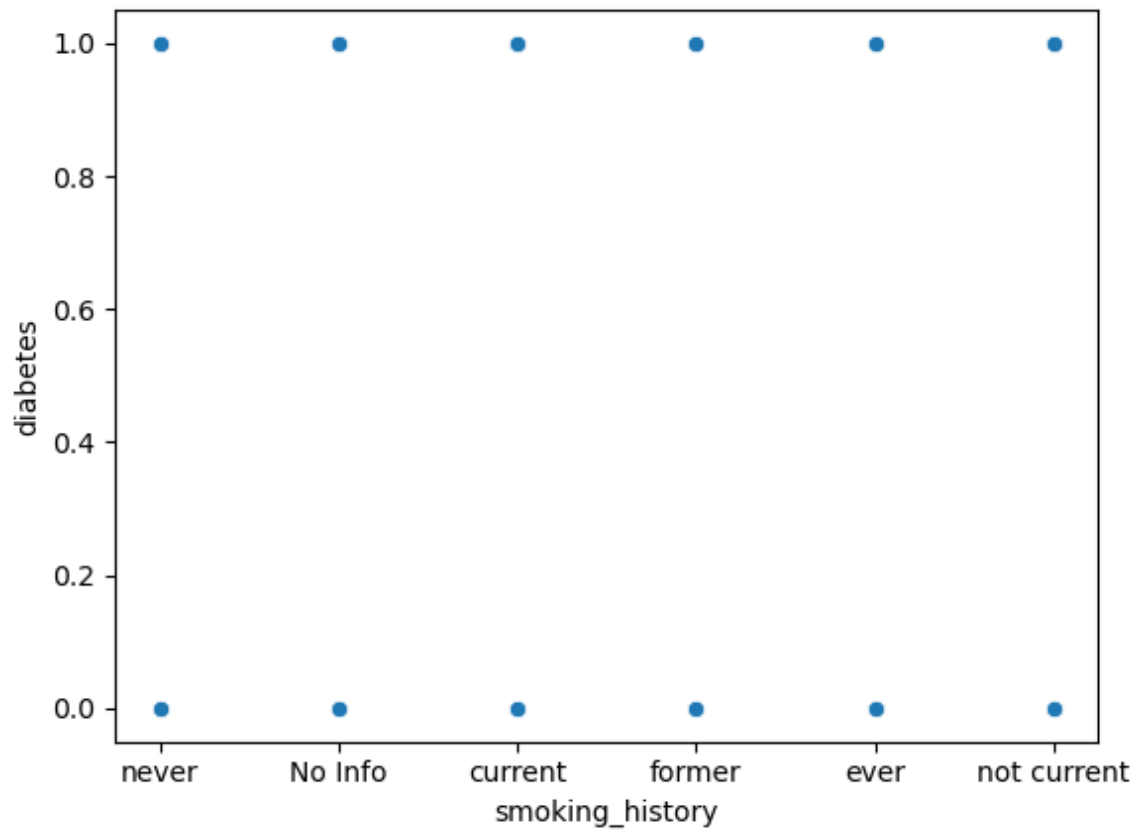




Now, lets check the correlation of the input and output Features.


```
In [13]: for i in df:  
          sns.scatterplot(df, y=df["diabetes"], x=df[i])  
          plt.show()
```











```
In [14]: df.columns
```

```
Out[14]: Index(['gender', 'age', 'hypertension', 'heart_disease', 'smoking_history',  
              'bmi', 'HbA1c_level', 'blood_glucose_level', 'diabetes'],  
             dtype='object')
```

```
In [15]: df["diabetes"].value_counts()
```

```
Out[15]: 0    91500  
         1     8500  
         Name: diabetes, dtype: int64
```

In [16]: `df[df.duplicated()]`

Out[16]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood
2756	Male	80.0	0	0	No Info	27.32	6.6	
3272	Female	80.0	0	0	No Info	27.32	3.5	
3418	Female	19.0	0	0	No Info	27.32	6.5	
3939	Female	78.0	1	0	former	27.32	3.5	
3960	Male	47.0	0	0	No Info	27.32	6.0	
...
99980	Female	52.0	0	0	never	27.32	6.1	
99985	Male	25.0	0	0	No Info	27.32	5.8	
99989	Female	26.0	0	0	No Info	27.32	5.0	
99990	Male	39.0	0	0	No Info	27.32	6.1	
99995	Female	80.0	0	0	No Info	27.32	6.2	

3854 rows × 9 columns



Here, we get the duplicate value but in this case we can't drop duplicate but all the details are important for ML.

Seperate data in X and Y as well as Split data into trainand Test

I am using a df1 data which was copy of the original data set.

In [17]: `x = df1.drop(["diabetes"], axis=1)`
`y = df1["diabetes"]`

In [18]: x

Out[18]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood
0	Female	80.0	0	1	never	25.19	6.6	
1	Female	54.0	0	0	No Info	27.32	6.6	
2	Male	28.0	0	0	never	27.32	5.7	
3	Female	36.0	0	0	current	23.45	5.0	
4	Male	76.0	1	1	current	20.14	4.8	
...
99995	Female	80.0	0	0	No Info	27.32	6.2	
99996	Female	2.0	0	0	No Info	17.37	6.5	
99997	Male	66.0	0	0	former	27.83	5.7	
99998	Female	24.0	0	0	never	35.42	4.0	
99999	Female	57.0	0	0	current	22.43	6.6	

100000 rows × 8 columns



```
In [19]: from sklearn.model_selection import train_test_split
train_x, test_x, train_y, test_y = train_test_split(x,y, random_state=50, t
est_size=0.25, stratify=y)
```

In [20]: train_x

Out[20]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood
28986	Female	32.0	0	0	never	30.47	6.2	
25188	Male	61.0	0	1	current	24.96	4.0	
37902	Male	30.0	0	0	ever	27.32	6.5	
55404	Female	61.0	0	0	ever	27.88	5.0	
69525	Female	15.0	0	0	No Info	27.32	6.1	
...
51401	Male	31.0	0	0	former	31.21	4.5	
6802	Male	16.0	0	0	No Info	27.32	4.8	
72319	Male	38.0	0	0	never	23.10	4.8	
31823	Female	50.0	0	0	current	19.87	4.0	
96394	Male	65.0	0	0	never	33.30	6.1	

75000 rows × 8 columns



Reset index

```
In [21]: train_x.reset_index(inplace=True, drop=True)
test_x.reset_index(inplace=True, drop=True)
train_y.reset_index(inplace=True, drop=True)
test_y.reset_index(inplace=True, drop=True)
```

```
In [22]: test_x
```

Out[22]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood
0	Female	53.0	0	0	No Info	27.32	6.0	
1	Female	16.0	0	0	No Info	27.32	6.6	
2	Female	31.0	0	0	No Info	27.32	5.7	
3	Female	38.0	0	0	current	39.01	4.0	
4	Female	59.0	0	0	never	27.32	6.5	
...
24995	Male	80.0	0	1	ever	28.24	8.8	
24996	Male	57.0	0	0	No Info	20.63	6.1	
24997	Male	37.0	0	0	never	27.32	4.8	
24998	Female	22.0	0	0	never	24.50	3.5	
24999	Female	63.0	0	0	No Info	35.68	9.0	

25000 rows × 8 columns



We can separate categorical and numerical features for encoding and scaling.

```
In [23]: train_cat = train_x.select_dtypes(include="object")
train_num = train_x.select_dtypes(include="number")
test_cat = test_x.select_dtypes(include="object")
test_num = test_x.select_dtypes(include="number")
```

```
In [24]: train_cat.head()
```

Out[24]:

	gender	smoking_history
0	Female	never
1	Male	current
2	Male	ever
3	Female	ever
4	Female	No Info

In [25]: `train_num.head()`

Out[25]:

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
0	32.0	0	0	30.47	6.2	126
1	61.0	0	1	24.96	4.0	130
2	30.0	0	0	27.32	6.5	158
3	61.0	0	0	27.88	5.0	130
4	15.0	0	0	27.32	6.1	100

In [26]: `test_cat.head()`

Out[26]:

	gender	smoking_history
0	Female	No Info
1	Female	No Info
2	Female	No Info
3	Female	current
4	Female	never

In [27]: `test_num.head()`

Out[27]:

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
0	53.0	0	0	27.32	6.0	160
1	16.0	0	0	27.32	6.6	80
2	31.0	0	0	27.32	5.7	160
3	38.0	0	0	39.01	4.0	200
4	59.0	0	0	27.32	6.5	145

Encoding By Using CatboostEncoder

Now for categorical encoding

In [28]: `import category_encoders as ce`

In [29]: `encoder = ce.CatBoostEncoder()
encoder.fit(train_cat, train_y)`

Out[29]:

▼	CatBoostEncoder
	CatBoostEncoder(cols=['gender', 'smoking_history'])

In [30]: `train_cat = encoder.transform(train_cat)
test_cat = encoder.transform(test_cat)`

In [31]: train_cat

Out[31]:

	gender	smoking_history
0	0.076283	0.095679
1	0.097311	0.099381
2	0.097311	0.119856
3	0.076283	0.119856
4	0.076283	0.040511
...
74995	0.097311	0.169112
74996	0.097311	0.040511
74997	0.097311	0.095679
74998	0.076283	0.099381
74999	0.097311	0.095679

75000 rows × 2 columns

Now concatenate train_cat and train_num as well as test_cat and test_num

In [32]: train_x = pd.concat([train_num, train_cat], axis=1)
test_x = pd.concat([test_num, test_cat], axis=1)

In [33]: train_x.head()

Out[33]:

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level	gender	smo
0	32.0	0	0	30.47	6.2	126	0.076283	
1	61.0	0	1	24.96	4.0	130	0.097311	
2	30.0	0	0	27.32	6.5	158	0.097311	
3	61.0	0	0	27.88	5.0	130	0.076283	
4	15.0	0	0	27.32	6.1	100	0.076283	

Now we can scaling the input variables

Sacling Features By Using MinMaxScaler

In [34]: from sklearn.preprocessing import StandardScaler, MinMaxScaler

```
In [35]: scaler = MinMaxScaler()
scaler.fit(train_x)
```

```
Out[35]:
▼ MinMaxScaler
MinMaxScaler()
```

```
In [36]: train_x = pd.DataFrame(scaler.transform(train_x), columns=train_x.columns)
test_x = pd.DataFrame(scaler.transform(test_x), columns=test_x.columns)
```

```
In [37]: test_x.head()
```

```
Out[37]:
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level	gender
0	0.662162	0.0	0.0	0.202031	0.454545	0.363636	0.76527
1	0.199199	0.0	0.0	0.202031	0.563636	0.000000	0.76527
2	0.386887	0.0	0.0	0.202031	0.400000	0.363636	0.76527
3	0.474474	0.0	0.0	0.338469	0.090909	0.545455	0.76527
4	0.737237	0.0	0.0	0.202031	0.545455	0.295455	0.76527

Basic Model Building And Evaluation

```
In [38]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
import xgboost as Xgb
from sklearn.naive_bayes import GaussianNB
```

```
In [39]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
classification_report, f1_score
```

```
In [40]: model1 = LogisticRegression(random_state=50)
model1.fit(train_x, train_y)
pred1 = model1.predict(test_x)
print("Testing Score :", recall_score(test_y, pred1))
print(classification_report(test_y, pred1))
print("Training Score :", model1.score(train_x, train_y))
```

```
Testing Score : 0.611764705882353
```

	precision	recall	f1-score	support
0	0.96	0.99	0.98	22875
1	0.88	0.61	0.72	2125
accuracy			0.96	25000
macro avg	0.92	0.80	0.85	25000
weighted avg	0.96	0.96	0.96	25000

```
Training Score : 0.9603866666666667
```

```
In [41]: model2 = KNeighborsClassifier()
model2.fit(train_x, train_y)
pred2 = model2.predict(test_x)
print("Testing Score :",recall_score(test_y, pred2))
print(classification_report(test_y, pred2))
print("Training Score :", model2.score(train_x, train_y))
```

Testing Score : 0.6202352941176471

	precision	recall	f1-score	support
0	0.97	0.99	0.98	22875
1	0.91	0.62	0.74	2125
accuracy			0.96	25000
macro avg	0.94	0.81	0.86	25000
weighted avg	0.96	0.96	0.96	25000

Training Score : 0.97064

```
In [42]: model3 = RandomForestClassifier(random_state=50)
model3.fit(train_x, train_y)
pred3 = model3.predict(test_x)
print("Testing Score :",recall_score(test_y, pred3))
print(classification_report(test_y, pred3))
print("Training Score :",model3.score(train_x, train_y))
```

Testing Score : 0.6781176470588235

	precision	recall	f1-score	support
0	0.97	1.00	0.98	22875
1	0.95	0.68	0.79	2125
accuracy			0.97	25000
macro avg	0.96	0.84	0.89	25000
weighted avg	0.97	0.97	0.97	25000

Training Score : 0.9991866666666667

```
In [43]: model4 = DecisionTreeClassifier()
model4.fit(train_x, train_y)
pred4 = model4.predict(test_x)
print("Testing Score :",recall_score(test_y, pred4))
print(classification_report(test_y, pred4))
print("Training Score :",model4.score(train_x, train_y))
```

Testing Score : 0.7369411764705882

	precision	recall	f1-score	support
0	0.98	0.97	0.97	22875
1	0.70	0.74	0.72	2125
accuracy			0.95	25000
macro avg	0.84	0.85	0.85	25000
weighted avg	0.95	0.95	0.95	25000

Training Score : 0.9992266666666667

```
In [44]: model5 = SVC(random_state=50)
model5.fit(train_x, train_y)
pred5 = model5.predict(test_x)
print("Testing Score :", recall_score(test_y, pred5))
print(classification_report(test_y, pred5))
print("Training Score :", model5.score(train_x, train_y))
```

Testing Score : 0.5609411764705883

	precision	recall	f1-score	support
0	0.96	1.00	0.98	22875
1	0.98	0.56	0.71	2125
accuracy			0.96	25000
macro avg	0.97	0.78	0.85	25000
weighted avg	0.96	0.96	0.96	25000

Training Score : 0.963

```
In [45]: model6 = Xgb.XGBClassifier()
model6.fit(train_x, train_y)
pred6 = model6.predict(test_x)
print("Testing Score :", recall_score(test_y, pred6))
print(classification_report(test_y, pred6))
print("Training Score :", model6.score(train_x, train_y))
```

Testing Score : 0.6856470588235294

	precision	recall	f1-score	support
0	0.97	1.00	0.98	22875
1	0.95	0.69	0.80	2125
accuracy			0.97	25000
macro avg	0.96	0.84	0.89	25000
weighted avg	0.97	0.97	0.97	25000

Training Score : 0.97644

```
In [46]: model7 = AdaBoostClassifier(random_state=50)
model7.fit(train_x, train_y)
pred7 = model7.predict(test_x)
print("Testing Score :", recall_score(test_y, pred7))
print(classification_report(test_y, pred7))
print("Training Score :", model7.score(train_x, train_y))
```

Testing Score : 0.6795294117647059

	precision	recall	f1-score	support
0	0.97	1.00	0.98	22875
1	0.97	0.68	0.80	2125
accuracy			0.97	25000
macro avg	0.97	0.84	0.89	25000
weighted avg	0.97	0.97	0.97	25000

Training Score : 0.9723866666666666

```
In [47]: model8 = GaussianNB()
model8.fit(train_x, train_y)
pred8 = model8.predict(test_x)
print("Testing Score :", recall_score(test_y, pred8))
print(classification_report(test_y, pred8))
print("Training Score :", model8.score(train_x, train_y))
```

Testing Score : 0.6357647058823529

	precision	recall	f1-score	support
0	0.96	0.93	0.95	22875
1	0.45	0.64	0.53	2125
accuracy			0.90	25000
macro avg	0.71	0.78	0.74	25000
weighted avg	0.92	0.90	0.91	25000

Training Score : 0.9034533333333333

Here, Decision Tree, Random forest and Xgboost model are overfitting of train and test dataset. So, we do a Hyper parameter tuning and Features selections.

HYPERPARAMETER TUNING

```
In [48]: #HYPERPERAMETER TUNING OF LOGISTIC REGRESSOR
from sklearn.model_selection import GridSearchCV
log = LogisticRegression()
params = { "tol" : [0.1,0.5,0.8,0.9], "C" : [1,2,8,6,9],
           "solver": ['lbfgs', "liblinear", "newton-cg", "newton-cholesky", "sag", "saga"]}
clf1 = GridSearchCV(log, params, cv=5, scoring="recall")
clf1.fit(train_x, train_y)
print(clf1.best_params_)
print(clf1.best_score_)
```

{'C': 8, 'solver': 'sag', 'tol': 0.9}
0.7007058823529412

```
In [49]: #HYPERPERAMETER TUNING OF NB
nb = GaussianNB()
params_nb = {'var_smoothing' : [0.96,0.25,0.30,0.40, 0.50]}
clf3 = GridSearchCV(nb, params_nb, cv=5, scoring="recall")
clf3.fit(train_x, train_y)
print(clf3.best_params_)
print(clf3.best_score_)
```

{'var_smoothing': 0.25}
0.5069803921568627

```
In [50]: #HYPERPERAMETER TUNING OF DECISION TREE
dt = DecisionTreeClassifier()
params_dt = {'criterion':['gini', 'entropy', 'log_loss'], 'max_depth' : [1,2
5,14,13,45,75,26], 'splitter':['best', 'random']}
clf5 = GridSearchCV(dt, params_dt, cv=5, scoring="recall")
clf5.fit(train_x, train_y)
print(clf5.best_params_)
print(clf5.best_score_)
```

```
{'criterion': 'gini', 'max_depth': 75, 'splitter': 'random'}
0.7408627450980392
```

```
In [51]: #HYPERPERAMETER TUNING OF RANDOMFOREST
rfc = RandomForestClassifier()
params_rfc = {"n_estimators" : [10,15,125,10,8,85], "max_depth" : [10,25,48,
85,42,3]}
clf6 = GridSearchCV(rfc, params_rfc, cv=5, scoring="recall")
clf6.fit(train_x, train_y)
print(clf6.best_params_)
print(clf6.best_score_)
```

```
{'max_depth': 85, 'n_estimators': 15}
0.6967843137254902
```

```
In [52]: #HYPERPERAMETER TUNING OF XGBOOST
xgb = Xgb.XGBClassifier()
params_xgb = {'eta': [0.1, 0.2, 0.3,0.4,0.5], 'n_estimators' : [10, 50, 10
0,12,15], 'max_depth': [3, 6, 9,14]}
clf7 = GridSearchCV(xgb, params_xgb, cv=5, scoring="recall")
clf7.fit(train_x, train_y)
print(clf7.best_params_)
print(clf7.best_score_)
```

```
{'eta': 0.5, 'max_depth': 14, 'n_estimators': 100}
0.7162352941176471
```

```
In [53]: #HYPERPERAMETER TUNING OF ADABOOST
adb = AdaBoostClassifier()
params_adb = {'n_estimators' : [10, 50, 100,12,15]}
clf8 = GridSearchCV(xgb, params_adb, cv=5, scoring="recall")
clf8.fit(train_x, train_y)
print(clf8.best_params_)
print(clf8.best_score_)
```

```
{'n_estimators': 100}
0.6986666666666667
```

```
In [54]: #best parameter for model
print("LogisticRegression score is :", clf1.best_params_)
print("GaussianNB score is :", clf3.best_params_)
print("DecisionTreeClassifier score is :", clf5.best_params_)
print("RandomForestClassifier score is :", clf6.best_params_)
print("XGB00ST score is :", clf7.best_params_)
print("AdaBoostClassifier score is :", clf8.best_params_)
```

```
LogisticRegression score is : {'C': 8, 'solver': 'sag', 'tol': 0.9}
GaussianNB score is : {'var_smoothing': 0.25}
DecisionTreeClassifier score is : {'criterion': 'gini', 'max_depth': 75,
'splitter': 'random'}
RandomForestClassifier score is : {'max_depth': 85, 'n_estimators': 15}
XGB00ST score is : {'eta': 0.5, 'max_depth': 14, 'n_estimators': 100}
AdaBoostClassifier score is : {'n_estimators': 100}
```

```
In [55]: #Score for all model
print("LogisticRegression score is :", clf1.best_score_)
print("GaussianNB score is :", clf3.best_score_)
print("DecisionTreeClassifier score is :", clf5.best_score_)
print("RandomForestClassifier score is :", clf6.best_score_)
print("XGB00ST score is :", clf7.best_score_)
print("AdaBoostClassifier score is :", clf8.best_score_)
```

```
LogisticRegression score is : 0.7007058823529412
GaussianNB score is : 0.5069803921568627
DecisionTreeClassifier score is : 0.7408627450980392
RandomForestClassifier score is : 0.6967843137254902
XGB00ST score is : 0.7162352941176471
AdaBoostClassifier score is : 0.6986666666666667
```

Feature Selection

```
In [56]: #Correlation
corr = train_x.corr()
corr.style.background_gradient(cmap='coolwarm')
```

Out[56]:

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
age	1.000000	0.251624	0.234324	0.336925	0.099437	0.110175
hypertension	0.251624	1.000000	0.115795	0.146456	0.082221	0.082860
heart_disease	0.234324	0.115795	1.000000	0.058238	0.067776	0.067703
bmi	0.336925	0.146456	0.058238	1.000000	0.083463	0.089174
HbA1c_level	0.099437	0.082221	0.067776	0.083463	1.000000	0.164899
blood_glucose_level	0.110175	0.082860	0.067703	0.089174	0.164899	1.000000
gender	-0.028546	0.014241	0.080006	-0.022681	0.018533	0.018533
smoking_history	0.322796	0.128296	0.094244	0.220150	0.053638	0.053638

```
In [57]: def correlation(dataset, threshold):
          col_corr = set()
          corr_matrix = dataset.corr()
          for i in range (len(corr_matrix.columns)):
              for j in range(i):
                  if abs(corr_matrix.iloc[i,j]) > threshold:
                      colname = corr_matrix.columns[i]
                      col_corr.add(colname)
          return col_corr
```

```
In [58]: corr_features = correlation(train_x, 0.7)
          len(set(corr_features))
```

Out[58]: 0

```
In [59]: corr_features
```

Out[59]: set()

```
In [60]: #Apply SelectKbest class to extract top Features
          from sklearn.feature_selection import SelectKBest, chi2
          bestfeatures = SelectKBest(score_func=chi2, k=7)
          fit = bestfeatures.fit(train_x, train_y)
```

```
In [61]: dfscores = pd.DataFrame(fit.scores_)
```

```
In [62]: dfcolumns = pd.DataFrame(x.columns)
```

```
In [63]: features = pd.concat([dfcolumns, dfscores], axis=1)
          features.columns = ["specs", "score"]
```

```
In [64]: features
```

Out[64]:

	specs	score
0	gender	766.343405
1	age	2699.337830
2	hypertension	2057.257217
3	heart_disease	102.219875
4	smoking_history	1243.196875
5	bmi	1696.135881
6	HbA1c_level	1.631301
7	blood_glucose_level	388.198706

```
In [65]: from sklearn.ensemble import ExtraTreesClassifier
          model = ExtraTreesClassifier()
          model.fit(train_x, train_y)
```


Out[65]:

▾ ExtraTreesClassifier
 ExtraTreesClassifier()

In [66]: `print(model.feature_importances_)`

```
[0.11904528 0.02451056 0.01631071 0.12639277 0.35885373 0.31576748
 0.00724419 0.03187528]
```

In [67]: `feat_importance = pd.Series(model.feature_importances_, index=x.columns)`
`feat_importance.nlargest(9).plot(kind="barh")`
`plt.show()`



In [68]: `fe_model = RandomForestClassifier(random_state=50)`
`fe_model.fit(train_x, train_y)`

Out[68]: `RandomForestClassifier`
`RandomForestClassifier(random_state=50)`

In [69]: `feature_scores = pd.Series(fe_model.feature_importances_, index=train_x.columns).sort_values(ascending=False)`

In [70]: `feature_scores`

```
Out[70]: HbA1c_level      0.388938
blood_glucose_level    0.327517
bmi                    0.122289
age                   0.100137
smoking_history        0.026238
hypertension           0.016730
heart_disease          0.010683
gender                 0.007469
dtype: float64
```

After the all Feature Selection method use then we decide to drop a gender, heart_disease column for best accuracy. So, start with second time generate Model

Start Whole Process of Model training in second time

```
In [71]: X = df1.drop(["gender", "heart_disease", "diabetes"], axis=1)
Y = df1["diabetes"]
```

```
In [72]: X
```

```
Out[72]:
```

	age	hypertension	smoking_history	bmi	HbA1c_level	blood_glucose_level
0	80.0	0	never	25.19	6.6	140
1	54.0	0	No Info	27.32	6.6	80
2	28.0	0	never	27.32	5.7	158
3	36.0	0	current	23.45	5.0	155
4	76.0	1	current	20.14	4.8	155
...
99995	80.0	0	No Info	27.32	6.2	90
99996	2.0	0	No Info	17.37	6.5	100
99997	66.0	0	former	27.83	5.7	155
99998	24.0	0	never	35.42	4.0	100
99999	57.0	0	current	22.43	6.6	90

100000 rows × 6 columns

```
In [73]: train_x1, test_x1, train_y1, test_y1 = train_test_split(X,Y, random_state=5
0,test_size=0.25, stratify=y)
```

In [74]: train_x1

Out[74]:

	age	hypertension	smoking_history	bmi	HbA1c_level	blood_glucose_level
28986	32.0	0	never	30.47	6.2	126
25188	61.0	0	current	24.96	4.0	130
37902	30.0	0	ever	27.32	6.5	158
55404	61.0	0	ever	27.88	5.0	130
69525	15.0	0	No Info	27.32	6.1	100
...
51401	31.0	0	former	31.21	4.5	160
6802	16.0	0	No Info	27.32	4.8	126
72319	38.0	0	never	23.10	4.8	100
31823	50.0	0	current	19.87	4.0	90
96394	65.0	0	never	33.30	6.1	130

75000 rows × 6 columns

```
In [75]: train_x1.reset_index(inplace=True, drop=True)
test_x1.reset_index(inplace=True, drop=True)
train_y1.reset_index(inplace=True, drop=True)
test_y1.reset_index(inplace=True, drop=True)
```

```
In [76]: train_cat = train_x1.select_dtypes(include="object")
train_num = train_x1.select_dtypes(include="number")
test_cat = test_x1.select_dtypes(include="object")
test_num = test_x1.select_dtypes(include="number")
```

In [77]: train_cat.head()

Out[77]:

	smoking_history
0	never
1	current
2	ever
3	ever
4	No Info

In [78]: encoder.fit(train_cat, train_y)

Out[78]:

```
▼ CatBoostEncoder
CatBoostEncoder(cols=['smoking_history'])
```

```
In [79]: train_cat = encoder.transform(train_cat)
test_cat = encoder.transform(test_cat)
```

In [80]: `test_cat.head()`

Out[80]:

	smoking_history
0	0.040511
1	0.040511
2	0.040511
3	0.099381
4	0.095679

In [81]: `train_x1 = pd.concat([train_num, train_cat], axis=1)`
`test_x1 = pd.concat([test_num, test_cat], axis=1)`

In [82]: `train_x1.head()`

Out[82]:

	age	hypertension	bmi	HbA1c_level	blood_glucose_level	smoking_history
0	32.0	0	30.47	6.2	126	0.095679
1	61.0	0	24.96	4.0	130	0.099381
2	30.0	0	27.32	6.5	158	0.119856
3	61.0	0	27.88	5.0	130	0.119856
4	15.0	0	27.32	6.1	100	0.040511

In [83]: `scaler.fit(train_x1)`

Out[83]:

▼ MinMaxScaler
 MinMaxScaler()

In [84]: `train_x1 = pd.DataFrame(scaler.transform(train_x1), columns=train_x1.columns)`
`test_x1 = pd.DataFrame(scaler.transform(test_x1), columns=test_x1.columns)`

In [85]: `train_x1.head()`

Out[85]:

	age	hypertension	bmi	HbA1c_level	blood_glucose_level	smoking_history
0	0.399399	0.0	0.238796	0.490909	0.209091	0.428981
1	0.762262	0.0	0.174486	0.090909	0.227273	0.457773
2	0.374374	0.0	0.202031	0.545455	0.354545	0.616985
3	0.762262	0.0	0.208567	0.272727	0.227273	0.616985
4	0.186687	0.0	0.202031	0.472727	0.090909	0.000000

```
In [86]: #LogisticRegression
log = LogisticRegression(C=1, solver="sag", tol=0.8)
log.fit(train_x1, train_y1)
pred1 = log.predict(test_x1)
print(classification_report(test_y1,pred1))
print(recall_score(test_y1, pred1))
```

	precision	recall	f1-score	support
0	0.97	0.98	0.98	22875
1	0.80	0.66	0.72	2125
accuracy			0.96	25000
macro avg	0.88	0.82	0.85	25000
weighted avg	0.95	0.96	0.95	25000

0.6574117647058824

```
In [87]: #GaussianNB
nb = GaussianNB(var_smoothing=0.25)
nb.fit(train_x1, train_y1)
pred3 = nb.predict(test_x1)
print(classification_report(test_y1,pred3))
print(recall_score(test_y1, pred3))
```

	precision	recall	f1-score	support
0	0.95	0.96	0.96	22875
1	0.52	0.49	0.51	2125
accuracy			0.92	25000
macro avg	0.74	0.73	0.73	25000
weighted avg	0.92	0.92	0.92	25000

0.49270588235294116

```
In [88]: #DecisionTreeClassifier
dt = DecisionTreeClassifier(criterion="entropy", max_depth=75, splitter="random")
dt.fit(train_x1, train_y1)
pred5 = dt.predict(test_x1)
print(classification_report(test_y1,pred5))
print(recall_score(test_y1, pred5))
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	22875
1	0.72	0.73	0.72	2125
accuracy			0.95	25000
macro avg	0.85	0.85	0.85	25000
weighted avg	0.95	0.95	0.95	25000

0.7312941176470589

```
In [89]: #RandomForestClassifier
rfc = RandomForestClassifier(max_depth=85 ,n_estimators= 15)
rfc.fit(train_x1, train_y1)
pred6 = rfc.predict(test_x1)
print(classification_report(test_y1,pred6))
print(recall_score(test_y1, pred6))
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	22875
1	0.89	0.69	0.77	2125
accuracy			0.97	25000
macro avg	0.93	0.84	0.88	25000
weighted avg	0.96	0.97	0.96	25000

0.6856470588235294

```
In [90]: #XGBClassifier
xgb = Xgb.XGBClassifier(eta=0.5 ,max_depth=14 ,n_estimators= 100)
xgb.fit(train_x1, train_y1)
pred7 = xgb.predict(test_x1)
print(classification_report(test_y1,pred7))
print(recall_score(test_y1, pred7))
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	22875
1	0.85	0.70	0.77	2125
accuracy			0.96	25000
macro avg	0.91	0.85	0.87	25000
weighted avg	0.96	0.96	0.96	25000

0.7021176470588235

```
In [91]: #AdaBoostClassifier
adb = AdaBoostClassifier(n_estimators= 100)
adb.fit(train_x1, train_y1)
pred8 = adb.predict(test_x1)
print(classification_report(test_y1,pred8))
print(recall_score(test_y1, pred8))
```

	precision	recall	f1-score	support
0	0.97	1.00	0.98	22875
1	0.98	0.68	0.80	2125
accuracy			0.97	25000
macro avg	0.98	0.84	0.89	25000
weighted avg	0.97	0.97	0.97	25000

0.676235294117647

```
In [92]: print('LogisticRegression score is ', recall_score(test_y1, pred1))
print('GaussianNB score is', recall_score(test_y1, pred3))
print('DecisionTreeClassifier score is', recall_score(test_y1, pred5))
print('RandomForestClassifier score is', recall_score(test_y1, pred6))
print('XGBClassifier score is', recall_score(test_y1, pred7))
print('AdaBoostClassifier score is', recall_score(test_y1, pred8))
```

```
LogisticRegression score is 0.6574117647058824
GaussianNB score is 0.49270588235294116
DecisionTreeClassifier score is 0.7312941176470589
RandomForestClassifier score is 0.6856470588235294
XGBClassifier score is 0.7021176470588235
AdaBoostClassifier score is 0.676235294117647
```

CONCLUSION :- IN ABOVE GENERATED MODEL IN RANDOM FOREST AND XGBOOST CLASSIFIER AND DECISIONTREE CLASSIFIER GIVE RECALL SCORE WAS LOW. ALSO MODEL PERFORMING OVERFITTING. SO THAT WE DO A OVERSAMPLING BECAUSE THE DATASET HAVE INBALANCED SO WE DO IT AND CHECK THE ACCURACY OF MODEL.

OVER SAMPLING

```
In [93]: xx = df2.drop(["gender", "heart_disease", "diabetes"], axis=1)
yy = df2["diabetes"]
```

```
In [94]: xx_cat = xx.select_dtypes(include="object")
xx_num = xx.select_dtypes(include="number")
```

```
In [95]: encoder.fit(xx_cat, yy)
xx_cat = encoder.transform(xx_cat)
```

```
In [96]: xx = pd.concat([xx_num, xx_cat], axis=1)
```

```
In [97]: from imblearn.over_sampling import SMOTE
os = SMOTE(random_state=50)
xos, yos = os.fit_resample(xx, yy)
```

```
In [98]: train_x11, test_x11, train_y11, test_y11 = train_test_split(xos, yos, random_state=50, test_size=0.2)
```

```
In [99]: train_x11.reset_index(inplace=True, drop=True)
train_y11.reset_index(inplace=True, drop=True)
test_x11.reset_index(inplace=True, drop=True)
test_y11.reset_index(inplace=True, drop=True)
```

In [100]: `train_x11.count()`

Out[100]:

age	146400
hypertension	146400
bmi	146400
HbA1c_level	146400
blood_glucose_level	146400
smoking_history	146400
dtype:	int64

In [101]: `test_x11.count()`

Out[101]:

age	36600
hypertension	36600
bmi	36600
HbA1c_level	36600
blood_glucose_level	36600
smoking_history	36600
dtype:	int64

In [102]: `scaler.fit(train_x11)`

Out[102]:

```

▼ MinMaxScaler
MinMaxScaler()

```

In [103]: `train_x11 = pd.DataFrame(scaler.transform(train_x11), columns=train_x11.columns)`
`test_x11 = pd.DataFrame(scaler.transform(test_x11), columns=test_x11.columns)`

In [104]: *#XGBClassifier*
`xgb.fit(train_x11, train_y11)`
`pred_1x = xgb.predict(test_x11)`
`print(classification_report(test_y11, pred_1x))`
`print("Testing Score :", recall_score(test_y11, pred_1x))`
`print("Training Score :", xgb.score(train_x11, train_y11))`


	precision	recall	f1-score	support
0	0.98	0.99	0.98	18205
1	0.99	0.98	0.98	18395
accuracy			0.98	36600
macro avg	0.98	0.98	0.98	36600
weighted avg	0.98	0.98	0.98	36600

Testing Score : 0.9751562924707801
 Training Score : 0.9985314207650273

In [105]: `from sklearn.metrics import ConfusionMatrixDisplay`


```
In [106]: print('XGBClassifier of confusion_matrix is:')
print(ConfusionMatrixDisplay.from_predictions(test_y11, pred_1x))
```

XGBClassifier of confusion_matrix is:
 <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x000001CA52307970>




```
In [107]: #RandomForestClassifier
rfc.fit(train_x11, train_y11)
pred_2r = rfc.predict(test_x11)
print(classification_report(test_y11, pred_2r))
print("Testing Score :", recall_score(test_y11, pred_2r))
print("Training Score :", rfc.score(train_x11, train_y11))
```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	18205
1	0.98	0.98	0.98	18395
accuracy			0.98	36600
macro avg	0.98	0.98	0.98	36600
weighted avg	0.98	0.98	0.98	36600

Testing Score : 0.9758086436531667
 Training Score : 0.9985792349726776

```
In [108]: print('RandomForestClassifier of confusion_matrix is:')
print(ConfusionMatrixDisplay.from_predictions(test_y11, pred_2r))
```

RandomForestClassifier of confusion_matrix is:
 <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x000001CA55304700>




```
In [109]: #DecisionTreeClassifier
dt.fit(train_x11, train_y11)
pred_2d = dt.predict(test_x11)
print(classification_report(test_y11,pred_2d))
print("Testing Score :",recall_score(test_y11, pred_2d))
print("Training Score :",rfc.score(train_x11, train_y11))
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	18205
1	0.97	0.97	0.97	18395
accuracy			0.97	36600
macro avg	0.97	0.97	0.97	36600
weighted avg	0.97	0.97	0.97	36600

Testing Score : 0.9709703723837999
 Training Score : 0.9985792349726776

```
In [110]: print('RandomForestClassifier of confusion_matrix is:')
          print(ConfusionMatrixDisplay.from_predictions(test_y11, pred_2d))
```

RandomForestClassifier of confusion_matrix is:
 <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x000001CA553C24A0>



Testing the new data for checking

```
In [111]: df["smoking_history"].unique()
```

```
Out[111]: array(['never', 'No Info', 'current', 'former', 'ever', 'not current'],
              dtype=object)
```

```
In [112]: xx.head()
```

```
Out[112]:
```

	age	hypertension	bmi	HbA1c_level	blood_glucose_level	smoking_history
0	80.0	0	25.19	6.6	140	0.095341
1	54.0	0	27.32	6.6	80	0.040598
2	28.0	0	27.32	5.7	158	0.095341
3	36.0	0	23.45	5.0	155	0.102087
4	76.0	1	20.14	4.8	155	0.102087

```
In [113]: new_df = {'age': 42, 'hypertension': 0, 'bmi': 25.36, 'HbA1c_level': 6.2, 'blood_glucose_level': 156, 'smoking_history': 'current'}
```

```
In [114]: index = [0]
```

```
In [115]: new_df = pd.DataFrame(new_df, index=index)
```

```
In [116]: new_df
```

```
Out[116]:
```

	age	hypertension	bmi	HbA1c_level	blood_glucose_level	smoking_history
0	42	0	25.36	6.2	156	current

```
In [117]: cat = new_df.select_dtypes(include="object")
num = new_df.select_dtypes(include="number")
```

```
In [118]: cat = encoder.transform(cat)
```

```
In [119]: cat
```

```
Out[119]:
```

	smoking_history
0	0.102087

```
In [120]: new_df = pd.concat([num, cat], axis=1)
```

```
In [121]: new_df
```

```
Out[121]:
```

	age	hypertension	bmi	HbA1c_level	blood_glucose_level	smoking_history
0	42	0	25.36	6.2	156	0.102087

```
In [122]: new_df = pd.DataFrame(scaler.transform(new_df), columns=new_df.columns)
```

```
In [123]: new_df
```

```
Out[123]:
```

	age	hypertension	bmi	HbA1c_level	blood_glucose_level	smoking_history
0	0.524525	0.0	0.179155	0.490909	0.345455	0.475153

```
In [124]: prediction = xgb.predict(new_df)
prediction
```

```
Out[124]: array([0])
```

```
In [125]: predic1 = rfc.predict(new_df)
predic1
```

```
Out[125]: array([0], dtype=int64)
```

```
In [126]: predic2 = dt.predict(new_df)
predic2
```

```
Out[126]: array([0], dtype=int64)
```

```
In [127]: new_df1= {'age': 86, 'hypertension': 1, 'bmi': 29.50, 'HbA1c_level':7.6 , 'blood_glucose_level': 17., 'smoking_history': 'never'}
```

```
In [128]: index = [0]
```

```
In [129]: new_df1 = pd.DataFrame(new_df1, index = index)
```

```
In [130]: new_df1
```

```
Out[130]:
```

	age	hypertension	bmi	HbA1c_level	blood_glucose_level	smoking_history
0	86	1	29.5	7.6	17.0	never

```
In [131]: cat = new_df1.select_dtypes(include="object")
num = new_df1.select_dtypes(include="number")
```

```
In [132]: cat = encoder.transform(cat)
```

```
In [133]: cat
```

```
Out[133]:
```

	smoking_history
0	0.095341

```
In [134]: new_df1 = pd.concat([num,cat], axis=1)
```

```
In [135]: new_df1
```

```
Out[135]:
```

	age	hypertension	bmi	HbA1c_level	blood_glucose_level	smoking_history
0	86	1	29.5	7.6	17.0	0.095341

```
In [136]: new_df1 = pd.DataFrame(scaler.transform(new_df1), columns=new_df1.columns)
```

```
In [137]: new_df1
```

```
Out[137]:
```

	age	hypertension	bmi	HbA1c_level	blood_glucose_level	smoking_history
0	1.075075	1.0	0.227474	0.745455	-0.286364	0.423021

```
In [138]: predic1 = xgb.predict(new_df1)
predic1
```

```
Out[138]: array([1])
```

```
In [139]: predic2 = rfc.predict(new_df1)
predic2
```

```
Out[139]: array([1], dtype=int64)
```

```
In [140]: predic3 = dt.predict(new_df1)
          predic3
```

```
Out[140]: array([0], dtype=int64)
```

CONCLUSION : The journey of developing machine learning models to solve our problem has been an iterative and comprehensive one, involving several crucial stages, including data processing, exploratory data analysis (EDA), basic model evaluation, hyperparameter tuning, and feature selection. Throughout this process, we employed nine different models to predict our target variable.

Upon evaluating these models, we observed the following performance scores:

Decision Tree Classifier: Testing Score: 0.9719, Training Score: 0.9986 Random Forest Classifier: Testing Score: 0.9748, Training Score: 0.9985 XGBoost Classifier: Testing Score: 0.9752, Training Score: 0.9985 Initially, our models did not meet our expectations in terms of testing scores, indicating a need for further improvement. In response, we decided to employ data sampling techniques to enhance model performance.

Among the three models, the XGBoost Classifier demonstrated the best performance, achieving the highest testing score of 0.9752. This result suggests that XGBoost was able to capture complex relationships within the data and make more accurate predictions compared to the other models.

In summary, our extensive efforts in data processing, EDA, model evaluation, hyperparameter tuning, and feature selection, combined with the use of advanced techniques like data sampling, culminated in the selection of the XGBoost Classifier as the most effective model for our task. This outcome underscores the importance of a systematic and iterative approach to machine learning model development, as it ultimately led to a model that meets or exceeds our desired level of performance.